

# Predicting Visitor Purchases with BigQuery ML

## Overview

[BigQuery ML](#) (BigQuery machine learning) is a feature in BigQuery where data analysts can create, train, evaluate, and predict with machine learning models with minimal coding. The Google Analytics Sample [Ecommerce dataset](#) that has millions of Google Analytics records for the [Google Merchandise Store](#) loaded into BigQuery. In this lab, you will use this data to run some typical queries that businesses would want to know about their customers' purchasing habits.

## Objectives

In this lab, you learn to perform the following tasks:

- Use BigQuery to find public datasets
- Query and explore the ecommerce dataset
- Create a training and evaluation dataset to be used for batch prediction
- Create a classification (logistic regression) model in BigQuery ML
- Evaluate the performance of your machine learning model
- Predict and rank the probability that a visitor will make a purchase

## Set up your environments

### Lab setup

For each lab, you get a new Google Cloud project and set of resources for a fixed time at no cost.

1. Sign in to Qwiklabs using an **incognito window**.
2. Note the lab's access time (for example, 1 : 15 : 00), and make sure you can finish within that time.  
There is no pause feature. You can restart if needed, but you have to start at the beginning.

3. When ready, click **Start lab**.
4. Note your lab credentials (**Username** and **Password**). You will use them to sign in to the Google Cloud Console.
5. Click **Open Google Console**.
6. Click **Use another account** and copy/paste credentials for **this** lab into the prompts.  
If you use other credentials, you'll receive errors or **incur charges**.
7. Accept the terms and skip the recovery resource page.

**Note:** Do not click **End Lab** unless you have finished the lab or want to restart it. This clears your work and removes the project.

## Open BigQuery Console

1. In the Google Cloud Console, select **Navigation menu** > **BigQuery**.  
The **Welcome to BigQuery in the Cloud Console** message box opens. This message box provides a link to the quickstart guide and lists UI updates.
2. Click **Done**.

## Access the course dataset

Once BigQuery is open, open the [data-to-insights](#) project in a new browser tab to bring this project into your BigQuery projects panel.

The field definitions for the **data-to-insights** ecommerce dataset are on the [\[UA\] BigQuery Export schema page](#). Keep the link open in a new tab for reference.

Filter Enter property name or value

<input type="checkbox"/>	Field name	Type	Mode	Key	Collation
<input type="checkbox"/>	<a href="#">visitorId</a>	INTEGER	NULLABLE		
<input type="checkbox"/>	<a href="#">visitNumber</a>	INTEGER	NULLABLE		
<input type="checkbox"/>	<a href="#">visitId</a>	INTEGER	NULLABLE		
<input type="checkbox"/>	<a href="#">visitStartTime</a>	INTEGER	NULLABLE		
<input type="checkbox"/>	<a href="#">date</a>	STRING	NULLABLE		
<input type="checkbox"/>	▶ <a href="#">totals</a>	RECORD	NULLABLE		
<input type="checkbox"/>	▶ <a href="#">trafficSource</a>	RECORD	NULLABLE		
<input type="checkbox"/>	▶ <a href="#">device</a>	RECORD	NULLABLE		
<input type="checkbox"/>	▶ <a href="#">geoNetwork</a>	RECORD	NULLABLE		
<input type="checkbox"/>	▶ <a href="#">customDimensions</a>	RECORD	REPEATED		
<input type="checkbox"/>	▶ <a href="#">hits</a>	RECORD	REPEATED		
<input type="checkbox"/>	<a href="#">fullVisitorId</a>	STRING	NULLABLE		
<input type="checkbox"/>	<a href="#">userId</a>	STRING	NULLABLE		
<input type="checkbox"/>	<a href="#">channelGrouping</a>	STRING	NULLABLE		
<input type="checkbox"/>	<a href="#">socialEngagementType</a>	STRING	NULLABLE		

## Task 1. Explore ecommerce data

**Scenario:** Your data analyst team exported the Google Analytics logs for an ecommerce website into BigQuery and created a new table of all the raw ecommerce visitor session data for you to explore. Using this data, you'll try to answer a few questions.

**Question:** Out of the total visitors who visited our website, what % made a purchase?

1. Click the query **EDITOR**.
2. Add the following to the New Query field:

```
#standardSQL
WITH visitors AS(
SELECT
COUNT(DISTINCT fullVisitorId) AS total_visitors
FROM `data-to-insights.ecommerce.web_analytics`
),
purchasers AS(
SELECT
COUNT(DISTINCT fullVisitorId) AS total_purchasers
FROM `data-to-insights.ecommerce.web_analytics`
WHERE totals.transactions IS NOT NULL
)
SELECT
    total_visitors,
    total_purchasers,
    total_purchasers / total_visitors AS conversion_rate
FROM visitors, purchasers
```

3. Click **Run**.

The result: 2.69%

Row	total_visitors	total_purchasers	conversion_rate
1	741721	20015	0.026984540008...

**Question:** What are the top 5 selling products?

4. Add the following query in the query **EDITOR**, and then click **Run**:

```
SELECT
    p.v2ProductName,
    p.v2ProductCategory,
    SUM(p.productQuantity) AS units_sold,
    ROUND(SUM(p.localProductRevenue/1000000),2) AS revenue
FROM `data-to-insights.ecommerce.web_analytics`,
UNNEST(hits) AS h,
UNNEST(h.product) AS p
GROUP BY 1, 2
ORDER BY revenue DESC
LIMIT 5;
```

The result:

Row	v2ProductName	v2ProductCategory	units_sold	revenue
1	Nest® Learning Thermostat 3rd Gen-USA - Stainless Steel	Nest-USA	17651	870976.95
2	Nest® Cam Outdoor Security Camera - USA	Nest-USA	16930	684034.55
3	Nest® Cam Indoor Security Camera - USA	Nest-USA	14155	548104.47
4	Nest® Protect Smoke + CO White Wired Alarm-USA	Nest-USA	6394	178937.6
5	Nest® Protect Smoke + CO White Battery Alarm-USA	Nest-USA	6340	178572.4

**Question:** How many visitors bought on subsequent visits to the website?

5. Run the following query to find out:

```
# visitors who bought on a return visit (could have bought on first as well
WITH all_visitor_stats AS (
SELECT
  fullvisitorid, # 741,721 unique visitors
  IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0, 1,
0) AS will_buy_on_return_visit
FROM `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid
)
SELECT
  COUNT(DISTINCT fullvisitorid) AS total_visitors,
  will_buy_on_return_visit
FROM all_visitor_stats
GROUP BY will_buy_on_return_visit
```

The results:

Row	total_visitors	will_buy_on_return_visit
1	729848	0
2	11873	1

Analyzing the results, you can see that  $(11873 / 729848) = 1.6\%$  of total visitors will return and purchase from the website. This includes the subset of visitors who bought on their very first session and then came back and bought again.

What are some of the reasons a typical ecommerce customer will browse but not buy until a later visit? Choose all that could apply.

The customer wants to comparison shop on other sites before making a purchase decision.

The customer is waiting for products to go on sale or other promotion

The customer is doing additional research

This behavior is very common for luxury goods where significant up-front research and comparison is required by the customer before deciding (think car purchases) but also true to a lesser extent for the merchandise on this site (t-shirts, accessories, etc).

In the world of online marketing, identifying and marketing to these future customers based on the characteristics of their first visit will increase conversion rates and reduce the outflow to competitor sites.

## Task 2. Select features and create your training dataset

Now you will create a Machine Learning model in BigQuery to predict whether or not a new user is likely to purchase in the future. Identifying these high-value users can help your marketing team target them with special promotions and ad campaigns.

Google Analytics captures a wide variety of dimensions and measures about a user's visit on this ecommerce website. Browse the complete list of fields in the [\[UA\] BigQuery Export schema Guide](#) and then [preview the demo dataset](#) to find useful features that will help a machine learning model understand the relationship between data about a visitor's first time on your website and whether they will return and make a purchase. Your team decides to test whether these two fields are good inputs for your classification model:

- `totals.bounces` (whether the visitor left the website immediately)
- `totals.timeOnSite` (how long the visitor was on our website)

What are the risks of only using the above two fields?

Whether a user bounces is highly correlated with their time on site (e.g. 0 seconds)

Only using time spent on the site ignores other potential useful columns (features)

Both of the above

Machine learning is only as good as the training data that is fed into it. If there isn't enough information for the model to determine and learn the relationship between your input features and your label (in this case, whether the visitor bought in the future) then you will not have an accurate model. While training a model on just these two fields is a start, you will see if they're good enough to produce an accurate model.

- In the query **EDITOR**, add the following query and then click **Run**:

```
SELECT
  * EXCEPT(fullVisitorId)
FROM
  # features
  (SELECT
    fullVisitorId,
    IFNULL(totals.bounces, 0) AS bounces,
    IFNULL(totals.timeOnSite, 0) AS time_on_site
```

```

FROM
  `data-to-insights.ecommerce.web_analytics`
WHERE
  totals.newVisits = 1)
JOIN
(SELECT
  fullvisitorid,
  IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0,
1, 0) AS will_buy_on_return_visit
FROM
  `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid)
USING (fullVisitorId)
ORDER BY time_on_site DESC
LIMIT 10;

```

Results:

Row	bounces	time_on_site	will_buy_on_return_visit
1	0	15047	0
2	0	12136	0
3	0	11201	0
4	0	10046	0
5	0	9974	0
6	0	9564	0



7	0	9520	0
8	0	9275	1
9	0	9138	0
10	0	8872	0

Which fields are the model features? What is the label (correct answer)?

The feature is will\_buy\_on\_return\_visit. The labels are bounces and time\_on\_site

The features are bounces and time\_on\_site. The label is will\_buy\_on\_return\_visit

The features are bounces and will\_buy\_on\_return\_visit. The label is time\_on\_site

Which fields are known after a visitor's first session? (Check all that apply)

time\_on\_site

will\_buy\_on\_return\_visit

visitId

bounces

Which field isn't known until later in the future after their first session?

time\_on\_site

visitId

will\_buy\_on\_return\_visit

bounces

**Discussion:** will\_buy\_on\_return\_visit is not known after the first visit. Again, you're predicting for a subset of users who returned to your website and purchased. Since you don't know the future at prediction time, you cannot say with certainty whether a new

visitor comes back and purchases. The value of building a ML model is to get the probability of future purchase based on the data gleaned about their first session.

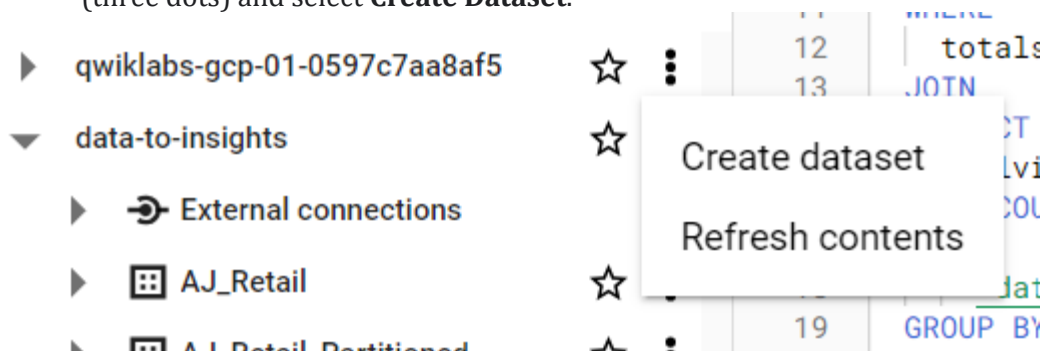
**Question:** Looking at the initial data results, do you think **time\_on\_site** and **bounces** will be a good indicator of whether the user will return and purchase or not?

**Answer:** It's often too early to tell before training and evaluating the model, but at first glance out of the top 10 **time\_on\_site**, only 1 customer returned to buy, which isn't very promising. Let's see how well the model does.

## Task 3. Create a BigQuery dataset to store models

Next, create a new BigQuery dataset which will also store your ML models.

1. In the left pane, click on your project name, and then click on the View action icon (three dots) and select **Create Dataset**.



2. In the **Create Dataset** dialog:
  - For **Dataset ID**, type **ecommerce**.
  - Leave the other values at their defaults.
3. Click **Create dataset**.

## Task 4. Select a BigQuery ML model type and specify options

Now that you have your initial features selected, you are now ready to create your first ML model in BigQuery.

There are the two model types to choose from:

Model	Model Type	Label Data type	Example
Forecasting	linear_reg	Numeric value (typically an integer or floating point)	Forecast sales figures for next year given historical sales data.
Classification	logistic_reg	0 or 1 for binary classification	Classify an email as spam or not spam given the context.

**Note:** There are many additional model types used in Machine Learning (like Neural Networks and decision trees) and available using libraries like [TensorFlow](#). At the time of writing, BigQuery ML supports the two listed above.

Which model type should you choose that will buy or won't buy?

Recommendation model (like matrix\_factorization etc.)

Classification model (like logistic\_reg etc.)

Forecasting model (like linear\_reg etc.)

1. Enter the following query to create a model and specify model options:

```
CREATE OR REPLACE MODEL `ecommerce.classification_model`  
OPTIONS  
(  
  model_type='logistic_reg',  
  labels = ['will_buy_on_return_visit']  
)  
AS  
#standardSQL  
SELECT  
  * EXCEPT(fullVisitorId)  
FROM  
  # features  
  (SELECT  
    fullVisitorId,  
    IFNULL(totals.bounces, 0) AS bounces,  
    IFNULL(totals.timeOnSite, 0) AS time_on_site  
  FROM  
    `data-to-insights.ecommerce.web_analytics`
```

```

WHERE
    totals.newVisits = 1
    AND date BETWEEN '20160801' AND '20170430') # train on first 9 months
JOIN
(SELECT
    fullvisitorid,
    IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0,
1, 0) AS will_buy_on_return_visit
FROM
    `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid)
USING (fullVisitorId)
;

```

2. Next, click **Run** to train your model.

Wait for the model to train (5 - 10 minutes).

**Note:** You cannot feed all of your available data to the model during training since you need to save some unseen data points for model evaluation and testing. To accomplish this, add a WHERE clause condition is being used to filter and train on only the first 9 months of session data in your 12 month dataset.

After your model is trained, you will see the message "This statement created a new model named qwiklabs-gcp-xxxxxxxxx:ecommerce.classification\_model".

3. Click **Go to model**.

Look inside the ecommerce dataset and confirm **classification\_model** now appears.

Next, you will evaluate the performance of the model against new unseen evaluation data.

Explorer

+ ADD

<

Type to search

?

Viewing resources.

SHOW STARRED ONLY

▼ qwiklabs-gcp-03-0822bed79b61

☆

⋮

▼ ecommerce

☆

⋮

▼ Models (1)

⋮

classification\_model

☆

⋮

SHOW MORE

SHOW MORE

SHOW MORE

► bigquery-public-data

★

⋮

SUMMARY

▼

Try new code-management features of RStudio in Preview

classification\_model

QUERY MODEL

DETAILS

TRAINING

EVALUATION

SCHEMA

Model type

LOGISTIC\_REGRESSION

Data location

US

Model Details

EDIT

Model ID

qwiklabs-gcp-03-0822bed79b61.ecommerce.classification\_model

Description

Date created

Nov 8, 2023, 9:29:45 AM UTC-8

Model expiration

Never

Date modified

Nov 8, 2023, 9:29:56 AM UTC-8

Data location

US

Model type

LOGISTIC\_REGRESSION

Loss type

Mean log loss

Training data

TEMPORARY TRAINING DATA TABLE

Evaluation data

TEMPORARY EVALUATION DATA TABLE

Training Options

Training options are the optional parameters that were added in the script to create this model.

Max allowed iterations

20

Actual iterations

10

L1 regularization

0.00

L2 regularization

0.00

Early stop

true

Min relative progress

0.01

Learn rate strategy

Line search

Line search initial learn rate

0.10

Calculate P Values

false


Data split method

Auto


View as

- ☒ Graphs  
☐ Table


Loss



Duration (seconds)



Learn rate



View as

- ☐ Graphs  
☒ Table

Iteration	Training Data Loss	Evaluation Data Loss	Learn Rate	Duration (seconds)
9	0.0727	0.0693	25.6	8.65
8	0.0730	0.0698	51.2	6.90
7	0.0755	0.0727	25.6	6.93
6	0.0838	0.0816	12.8	8.39
5	0.1083	0.1066	6.4	8.16
4	0.1709	0.1697	3.2	7.34
3	0.2919	0.2912	1.6	6.83
2	0.4419	0.4415	0.8	8.68
1	0.5665	0.5663	0.4	8.60
0	0.6473	0.6472	0.2	4.89

classification\_model

DETAILS TRAINING EVALUATION SCHEMA


Aggregate Metrics ?

Threshold ?	0.5000
Precision ?	0.0000
Recall ?	0.0000
Accuracy ?	0.9847
F1 score ?	0.0000
Log loss ?	0.0693
ROC AUC ?	0.7856

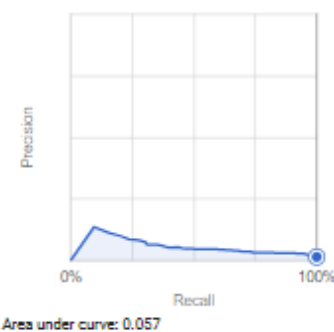
Score threshold

Positive class threshold ?	<input type="range" value="0.0070"/>
Positive class	1
Negative class	0
Precision ?	0.0148
Recall ?	1.0000
Accuracy ?	0.0148
F1 score ?	0.0292


Precision-recall by threshold ?



Precision-recall curve ?



ROC curve ?



Confusion matrix

Item counts

This table shows how often the model classified each label correctly (in blue), and which labels were most often confused for that label (in gray).

True label	Predicted label	
	1	0
1	100%	0%
0	100%	0%

classification\_model

QUERY M

DETAILS

TRAINING

EVALUATION

SCHEMA

Labels

Filter

Enter property name or value

Field name	Type	Mode	Description
<a href="#">predicted_will_buy_on_return_visit</a>	INT64	NULLABLE	

Features

Filter

Enter property name or value

Field name	Type	Mode	Description
<a href="#">bounces</a>	INT64	NULLABLE	
<a href="#">time_on_site</a>	INT64	NULLABLE	

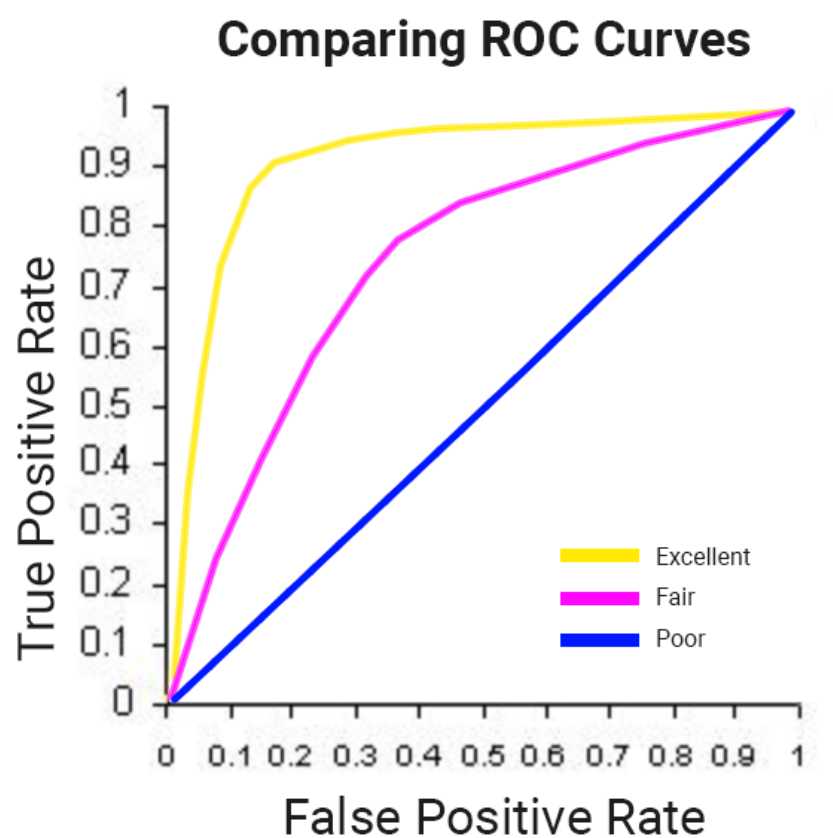
## Task 5. Evaluate classification model performance

### Select your performance criteria

For classification problems in ML, you want to minimize the False Positive Rate (predict that the user will return and purchase and they don't) and maximize the True Positive Rate (predict that the user will return and purchase and they do).

This relationship is visualized with a ROC (Receiver Operating Characteristic) curve like the one shown here, where you try to maximize the area under the curve or AUC:





In BigQuery ML, **roc\_auc** is simply a queryable field when evaluating your trained ML model.

- Now that training is complete, you can evaluate how well the model performs by running this query using `ML.EVALUATE`:

```
SELECT
  roc_auc,
  CASE
    WHEN roc_auc > .9 THEN 'good'
    WHEN roc_auc > .8 THEN 'fair'
    WHEN roc_auc > .7 THEN 'not great'
    ELSE 'poor' END AS model_quality
FROM
  ML.EVALUATE(MODEL ecommerce.classification_model, (
SELECT
  * EXCEPT(fullVisitorId)
FROM
  # features
  (SELECT
```

```

    fullVisitorId,
    IFNULL(totals.bounces, 0) AS bounces,
    IFNULL(totals.timeOnSite, 0) AS time_on_site
FROM
  `data-to-insights.ecommerce.web_analytics`
WHERE
  totals.newVisits = 1
  AND date BETWEEN '20170501' AND '20170630') # eval on 2 months
JOIN
  (SELECT
    fullvisitorid,
    IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0,
1, 0) AS will_buy_on_return_visit
  FROM
    `data-to-insights.ecommerce.web_analytics`
  GROUP BY fullvisitorid)
  USING (fullVisitorId)
));

```

You should see the following result:

Row	roc_auc	model_quality
1	0.724588	not great

After evaluating your model you get a **roc\_auc** of 0.72, which shows that the model has not great predictive power. Since the goal is to get the area under the curve as close to 1.0 as possible, there is room for improvement.

## Task 6. Improve model performance with feature engineering

As was hinted at earlier, there are many more features in the dataset that may help the model better understand the relationship between a visitor's first session and the likelihood that they will purchase on a subsequent visit.

Add some new features and create a second machine learning model called `classification_model_2`:

- How far the visitor got in the checkout process on their first visit
- Where the visitor came from (traffic source: organic search, referring site etc.)
- Device category (mobile, tablet, desktop)
- Geographic information (country)

1. Create this second model by running the below query:

```
CREATE OR REPLACE MODEL `ecommerce.classification_model_2`
OPTIONS
  (model_type='logistic_reg', labels = ['will_buy_on_return_visit']) AS
WITH all_visitor_stats AS (
SELECT
  fullvisitorid,
  IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0, 1,
0) AS will_buy_on_return_visit
FROM `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid
)
# add in new features
SELECT * EXCEPT(unique_session_id) FROM (
  SELECT
    CONCAT(fullvisitorid, CAST(visitId AS STRING)) AS unique_session_id,
    # labels
    will_buy_on_return_visit,
    MAX(CAST(h.eCommerceAction.action_type AS INT64)) AS
latest_ecommerce_progress,
    # behavior on the site
    IFNULL(totals.bounces, 0) AS bounces,
    IFNULL(totals.timeOnSite, 0) AS time_on_site,
    totals.pageviews,
    # where the visitor came from
    trafficSource.source,
    trafficSource.medium,
```

```

        channelGrouping,
        # mobile or desktop
        device.deviceCategory,
        # geographic
        IFNULL(geoNetwork.country, "") AS country
FROM `data-to-insights.ecommerce.web_analytics`,
    UNNEST(hits) AS h
JOIN all_visitor_stats USING(fullvisitorid)
WHERE 1=1
    # only predict for new visits
    AND totals.newVisits = 1
    AND date BETWEEN '20160801' AND '20170430' # train 9 months
GROUP BY
unique_session_id,
will_buy_on_return_visit,
bounces,
time_on_site,
totals.pageviews,
trafficSource.source,
trafficSource.medium,
channelGrouping,
device.deviceCategory,
country
);

```

**Note:** You are still training on the same first 9 months of data, even with this new model. It's important to have the same training dataset so you can be certain a better model output is attributable to better input features and not new or different training data.

A key new feature that was added to the training dataset query is the maximum checkout progress each visitor reached in their session, which is recorded in the field `hits.eCommerceAction.action_type`. If you search for that field in the [field definitions](#) you will see the field mapping of 6 = Completed Purchase.

As an aside, the web analytics dataset has nested and repeated fields like [ARRAYS](#) which need to be broken apart into separate rows in your dataset. This is accomplished by using the `UNNEST()` function, which you can see in the above query.

Wait for the new model to finish training (5-10 minutes).


Aggregate Metrics ?

Threshold ?	0.5000
Precision ?	0.2000
Recall ?	0.0263
Accuracy ?	0.9834
F1 score ?	0.0465
Log loss ?	0.0582
ROC AUC ?	0.9221

Score threshold

Positive class threshold ?	<input type="range" value="0.0021"/>
Positive class	1
Negative class	0
Precision ?	0.0154
Recall ?	1.0000
Accuracy ?	0.0154
F1 score ?	0.0302

Precision-recall by threshold ?



Precision-recall curve ?



ROC curve ?



Confusion matrix

☐ Item counts

This table shows how often the model classified each label correctly (in blue), and which labels were most often confused for that label (in gray).

True label	Predicted label	
	1	0
1	100%	0%
0	100%	0%

## classification\_model\_2

[QUERY](#)

DETAILS TRAINING EVALUATION **SCHEMA**

Field name	Type	Mode	Description
<a href="#">predicted_will_buy_on_return_visit</a>	INT64	NULLABLE	

## Features

 **Filter** Enter property name or value

Field name	Type	Mode	Description
<a href="#">latest_ecommerce_progress</a>	INT64	NULLABLE	
<a href="#">bounces</a>	INT64	NULLABLE	
<a href="#">time_on_site</a>	INT64	NULLABLE	
<a href="#">pageviews</a>	INT64	NULLABLE	
<a href="#">source</a>	STRING	NULLABLE	
<a href="#">medium</a>	STRING	NULLABLE	
<a href="#">channelGrouping</a>	STRING	NULLABLE	
<a href="#">deviceCategory</a>	STRING	NULLABLE	
<a href="#">country</a>	STRING	NULLABLE	

- Evaluate this new model to see if there is better predictive power by running the below query:

```
#standardSQL
```

```
SELECT
  roc_auc,
  CASE
    WHEN roc_auc > .9 THEN 'good'
    WHEN roc_auc > .8 THEN 'fair'
    WHEN roc_auc > .7 THEN 'not great'
    ELSE 'poor' END AS model_quality
FROM
  ML.EVALUATE(MODEL ecommerce.classification_model_2, (
WITH all_visitor_stats AS (
SELECT
  fullvisitorid,
  IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0, 1,
0) AS will_buy_on_return_visit
FROM `data-to-insights.ecommerce.web_analytics`
```

```

    GROUP BY fullvisitorid
)
# add in new features
SELECT * EXCEPT(unique_session_id) FROM (
    SELECT
        CONCAT(fullvisitorid, CAST(visitId AS STRING)) AS unique_session_id,
        # labels
        will_buy_on_return_visit,
        MAX(CAST(h.eCommerceAction.action_type AS INT64)) AS
latest_ecommerce_progress,
        # behavior on the site
        IFNULL(totals.bounces, 0) AS bounces,
        IFNULL(totals.timeOnSite, 0) AS time_on_site,
        totals.pageviews,
        # where the visitor came from
        trafficSource.source,
        trafficSource.medium,
        channelGrouping,
        # mobile or desktop
        device.deviceCategory,
        # geographic
        IFNULL(geoNetwork.country, "") AS country
    FROM `data-to-insights.ecommerce.web_analytics`,
        UNNEST(hits) AS h
    JOIN all_visitor_stats USING(fullvisitorid)
    WHERE 1=1
        # only predict for new visits
        AND totals.newVisits = 1
        AND date BETWEEN '20170501' AND '20170630' # eval 2 months
    GROUP BY
        unique_session_id,
        will_buy_on_return_visit,
        bounces,
        time_on_site,
        totals.pageviews,
        trafficSource.source,
        trafficSource.medium,
        channelGrouping,
        device.deviceCategory,
        country
)
));

```

(Output)

Row	roc_auc	model_quality
1	0.910382	good

With this new model you now get a **roc\_auc** of 0.91 which is significantly better than the first model.

Now that you have a trained model, time to make some predictions.

## Task 7. Predict which new visitors will come back and purchase

Next you will write a query to predict which new visitors will come back and make a purchase.

- Run the prediction query below which uses the improved classification model to predict the probability that a first-time visitor to the Google Merchandise Store will make a purchase in a later visit:

```
SELECT
*
FROM
  ml.PREDICT(MODEL `ecommerce.classification_model_2`,
  (
WITH all_visitor_stats AS (
SELECT
  fullvisitorid,
  IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0, 1,
0) AS will_buy_on_return_visit
FROM `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid
)
SELECT
  CONCAT(fullvisitorid, '-',CAST(visitId AS STRING)) AS
unique_session_id,
  # labels
  will_buy_on_return_visit,
```



```

        MAX(CAST(h.eCommerceAction.action_type AS INT64)) AS
latest_ecommerce_progress,
    # behavior on the site
    IFNULL(totals.bounces, 0) AS bounces,
    IFNULL(totals.timeOnSite, 0) AS time_on_site,
    totals.pageviews,
    # where the visitor came from
    trafficSource.source,
    trafficSource.medium,
    channelGrouping,
    # mobile or desktop
    device.deviceCategory,
    # geographic
    IFNULL(geoNetwork.country, "") AS country
FROM `data-to-insights.ecommerce.web_analytics`,
    UNNEST(hits) AS h
JOIN all_visitor_stats USING(fullvisitorid)
WHERE
    # only predict for new visits
    totals.newVisits = 1
    AND date BETWEEN '20170701' AND '20170801' # test 1 month
GROUP BY
    unique_session_id,
    will_buy_on_return_visit,
    bounces,
    time_on_site,
    totals.pageviews,
    trafficSource.source,
    trafficSource.medium,
    channelGrouping,
    device.deviceCategory,
    country
)
)
ORDER BY
    predicted_will_buy_on_return_visit DESC;

```

The predictions are made in the last 1 month (out of 12 months) of the dataset.

Your model will now output the predictions it has for those July 2017 ecommerce sessions. You can see three newly added fields:

- predicted\_will\_buy\_on\_return\_visit: whether the model thinks the visitor will buy later (1 = yes)
- predicted\_will\_buy\_on\_return\_visit\_probs.label: the binary classifier for yes / no

- `predicted_will_buy_on_return_visit_probs.prob`: the confidence the model has in it's prediction (1 = 100%)

Query results

SAVE RESULTS

EXPLORE DATA

JOB INFORMATION		RESULTS	CHART	PREVIEW	JSON	EXECUTION DETAILS		EXECUTION GRAPH									
Row	predicted_will_buy	predict_label	predict_prob	unique_session_id	will_buy_on_return	latest_ecommerce	bounces	time_on_site	pageviews	source	medium	channelGrouping	deviceCategory	country			
1	1	1	0.5755152673300	7344038342499021225-15011...	0	6	0	2515	61	mail.googleplex.com	referral	Referral	desktop	United States			
2	1	1	0.4244847326600	0070425995371344990-14998...	0	6	0	573	14	sites.google.com	referral	Referral	desktop	United States			
3	1	1	0.5959245399600	476662364150402239-15003...	1	6	0	1305	19	gdata.googleplex.com	referral	Referral	desktop	United States			
4	1	1	0.4604175460300	58070922137...	0	6	0	331	10	sites.google.com	referral	Referral	desktop	United States			
5	1	1	0.5355235522600	7774240357627896789-14994...	0	6	0	608	11	gdata.googleplex.com	referral	Referral	desktop	United States			
6	1	1	0.4644764477300	7019345936989504727-14999...	0	6	0	595	11	gdata.googleplex.com	referral	Referral	desktop	United States			
7	1	1	0.5230612058300	1849247637034984738-15006...	0	6	0	243	11	sites.google.com	referral	Referral	desktop	United States			
8	1	1	0.4769387941600	52246572493...	0	6	0										
9	1	1	0.5224657249300	47753427506...	0	6	0										
10	1	1	0.5347817355600	1413926032042347905-15007...	0	6	0										

## Results

- Of the top 6% of first-time visitors (sorted in decreasing order of predicted probability), more than 6% make a purchase in a later visit.
- These users represent nearly 50% of all first-time visitors who make a purchase in a later visit.
- Overall, only 0.7% of first-time visitors make a purchase in a later visit.
- Targeting the top 6% of first-time increases marketing ROI by 9x vs targeting them all!

## Additional information

**roc\_auc** is just one of the performance metrics available during model evaluation. Also available are [accuracy, precision, and recall](#). Knowing which performance metric to rely on is highly dependent on what your overall objective or goal is.

## Congratulations!

You created a machine learning model using just SQL.

# Challenge

## Summary

In the previous two tasks you saw the power of feature engineering at work in improving our models performance. However, we still may be able to improve our performance by exploring other model types. For classification problems, BigQuery ML also supports the following model types:

- [Deep Neural Networks](#).
- [Boosted Decision Trees \(XGBoost\)](#).
- [AutoML Tables Models](#).
- [Importing Custom TensorFlow Models](#).

## Task

Though our linear classification (logistic regression) model performed well after feature engineering, it may be too simple of a model to fully capture the relationship between the features and the label. Using the same dataset and labels as you did in Task 6 to create the model `ecommerce.classification_model_2`, your challenge is to create a XGBoost Classifier.

**Note: Hint :** Use following options for `Boosted_Tree_Classifier`:

1. `L2_reg = 0.1`
2. `num_parallel_tree = 8`
3. `max_tree_depth = 10`

You may need to look at the documentation linked above to see the exact syntax. The model will take around 7 minutes to train. The solution can be found in the solution section below if you need help writing the query. pa

## Solution:

This is the solution that you require in order to create a XGBoost Classifier:

```

CREATE OR REPLACE MODEL `ecommerce.classification_model_3`
  OPTIONS
    (model_type='BOOSTED_TREE_CLASSIFIER' , l2_reg = 0.1,
num_parallel_tree = 8, max_tree_depth = 10,
    labels = ['will_buy_on_return_visit']) AS
  WITH all_visitor_stats AS (
  SELECT
    fullvisitorid,
    IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS
NULL) > 0, 1, 0) AS will_buy_on_return_visit
  FROM `data-to-insights.ecommerce.web_analytics`
  GROUP BY fullvisitorid
  )
  # add in new features
  SELECT * EXCEPT(unique_session_id) FROM (
  SELECT
    CONCAT(fullvisitorid, CAST(visitId AS STRING)) AS
unique_session_id,
    # labels
    will_buy_on_return_visit,
    MAX(CAST(h.eCommerceAction.action_type AS INT64)) AS
latest_ecommerce_progress,
    # behavior on the site
    IFNULL(totals.bounces, 0) AS bounces,
    IFNULL(totals.timeOnSite, 0) AS time_on_site,
    totals.pageviews,
    # where the visitor came from
    trafficSource.source,
    trafficSource.medium,
    channelGrouping,
    # mobile or desktop
    device.deviceCategory,
    # geographic
    IFNULL(geoNetwork.country, "") AS country
  FROM `data-to-insights.ecommerce.web_analytics`,
    UNNEST(hits) AS h
  JOIN all_visitor_stats USING(fullvisitorid)
  WHERE 1=1
    # only predict for new visits
    AND totals.newVisits = 1
    AND date BETWEEN '20160801' AND '20170430' # train 9 months
  GROUP BY
    unique_session_id,
    will_buy_on_return_visit,
    bounces,
    time_on_site,

```

```

        totals.pageviews,
        trafficSource.source,
        trafficSource.medium,
        channelGrouping,
        device.deviceCategory,
        country
    );

```

Let us now evaluate our model and see how we did:

#standardSQL

```

SELECT
    roc_auc,
    CASE
        WHEN roc_auc > .9 THEN 'good'
        WHEN roc_auc > .8 THEN 'fair'
        WHEN roc_auc > .7 THEN 'not great'
        ELSE 'poor' END AS model_quality
FROM
    ML.EVALUATE(MODEL ecommerce.classification_model_3, (
WITH all_visitor_stats AS (
SELECT
    fullvisitorid,
    IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0, 1,
0) AS will_buy_on_return_visit
FROM `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid
)
# add in new features
SELECT * EXCEPT(unique_session_id) FROM (
    SELECT
        CONCAT(fullvisitorid, CAST(visitId AS STRING)) AS unique_session_id,
        # labels
        will_buy_on_return_visit,
        MAX(CAST(h.eCommerceAction.action_type AS INT64)) AS
latest_ecommerce_progress,
        # behavior on the site
        IFNULL(totals.bounces, 0) AS bounces,
        IFNULL(totals.timeOnSite, 0) AS time_on_site,
        totals.pageviews,
        # where the visitor came from
        trafficSource.source,
        trafficSource.medium,
        channelGrouping,
        # mobile or desktop

```

```

        device.deviceCategory,
        # geographic
        IFNULL(geoNetwork.country, "") AS country
FROM `data-to-insights.ecommerce.web_analytics`,
    UNNEST(hits) AS h
JOIN all_visitor_stats USING(fullvisitorid)
WHERE 1=1
    # only predict for new visits
    AND totals.newVisits = 1
    AND date BETWEEN '20170501' AND '20170630' # eval 2 months
GROUP BY
    unique_session_id,
    will_buy_on_return_visit,
    bounces,
    time_on_site,
    totals.pageviews,
    trafficSource.source,
    trafficSource.medium,
    channelGrouping,
    device.deviceCategory,
    country
)
));

```

Our roc\_auc has increased by about .02 to around .94!

Row	roc_auc	model_quality
1	0.921795204795...	good

**Note :** Your exact values will differ due to the randomness involved in the training process.

It's a small change in the roc\_auc, but note that since 1 is a perfect roc\_auc, it gets more difficult to improve the metric the closer to 1 it gets.

This is a great example of how easy it is in BigQuery ML to try out different model types with different options to see how they perform. We were able to use a much more complex model type by only changing one line of SQL.

One may reasonably ask “Where did the choices for these options come from?”, and the answer is experimentation! When you are trying to find the best model type for your problems, then one has to experiment with different sets of options in a process known as hyperparameter tuning.

Let's finish up by generating predictions with our improved model and see how they compare to those we generated before. By using a **Boosted tree classifier model**, you can observe a slight improvement of 0.2 in our ROC AUC compared to the previous model. The query below will predict which new visitors will come back and make a purchase:

```
SELECT
*
FROM
  ml.PREDICT(MODEL `ecommerce.classification_model_3`,
  (
WITH all_visitor_stats AS (
SELECT
  fullvisitorid,
  IF(COUNTIF(totals.transactions > 0 AND totals.newVisits IS NULL) > 0, 1,
0) AS will_buy_on_return_visit
FROM `data-to-insights.ecommerce.web_analytics`
GROUP BY fullvisitorid
)
SELECT
  CONCAT(fullvisitorid, '-',CAST(visitId AS STRING)) AS
unique_session_id,
  # labels
  will_buy_on_return_visit,
  MAX(CAST(h.eCommerceAction.action_type AS INT64)) AS
latest_ecommerce_progress,
  # behavior on the site
  IFNULL(totals.bounces, 0) AS bounces,
  IFNULL(totals.timeOnSite, 0) AS time_on_site,
  totals.pageviews,
  # where the visitor came from
  trafficSource.source,
  trafficSource.medium,
  channelGrouping,
  # mobile or desktop
  device.deviceCategory,
  # geographic
  IFNULL(geoNetwork.country, "") AS country
FROM `data-to-insights.ecommerce.web_analytics`,
  UNNEST(hits) AS h
JOIN all_visitor_stats USING(fullvisitorid)
WHERE
  # only predict for new visits
  totals.newVisits = 1
  AND date BETWEEN '20170701' AND '20170801' # test 1 month
GROUP BY
  unique_session_id,
```

```
will_buy_on_return_visit,  
bounces,  
time_on_site,  
totals.pageviews,  
trafficSource.source,  
trafficSource.medium,  
channelGrouping,  
device.deviceCategory,  
country  
)  
)  
ORDER BY  
  predicted_will_buy_on_return_visit DESC;  
Copied!  
content_copy
```

The output now shows a classification model that can better predict the probability that a first-time visitor to the Google Merchandise Store will make a purchase in a later visit. By comparing the result above with the previous model shown in Task 7, you can see the confidence the model has in its predictions is more accurate when compared to the logistic\_regression model type.

**Save results to “Boosted tree classifier model results.csv”**

## End your lab

When you have completed your lab, click **End Lab**. Google Cloud Skills Boost removes the resources you’ve used and cleans the account for you.

You will be given an opportunity to rate the lab experience. Select the applicable number of stars, type a comment, and then click **Submit**.